

Four minutes, **ten slides.**

What to say on every slide, **29 September.**

[← all documents](#) · Kat Titterton · San Francisco · 29 September 2026

Run of show · 29 Sept

Four minutes, ten slides, about 530 spoken words. The spine: **self, us, now** (02, 03), then **why, how, what** (04, 05, 06), then **who else, where it goes, the ask** (07–10). 07 and 08 are the fastest stretch: one beat each.

SLIDE	CLOCK	SEC	CLICKS	BEAT	ONE LINE TO LAND
01 Control is an illusion	0:00–0:12	12	1	hook	Confidence without control.
02 Founder	0:12–0:40	28	1	self	This is the thing I kept needing and nobody sold.
03 Do my cells look normal?	0:40–1:05	25	1	us, now	Cells are the input nobody certifies.
04 One lab, or six?	1:05–1:45	40	3	why now	The model isn't the hard part. The reference is.
05 It passed every check	1:45–2:20	35	3	how	Your own past is not a reference. The field is.
06 The product	2:20–2:45	25	1	what	Four months becomes one day.
07 Identity, yes. Behaviour, no.	2:45–2:55	10	1	who else	The other half of the certificate.

SLIDE	CLOCK	SEC	CLICKS	BEAT	ONE LINE TO LAND
08 Three buyers, one certificate	2:55–3:10	15	1	who buys	The labs that build the reference use it free.
09 The ugliest dataset in biology	3:10–3:35	25	1	the build	Name the failure, not just flag it.
10 Every plate, placed in the field	3:35–4:00	25	0	vision, ask	Have images of U2OS cells? Talk to me.

Captions stay hidden while you present (Tweaks › captions); the shared link and PDF show them. If you're late at 2:20, cut the "why now" sentence on 04 and the "who pays" line on 06.

01 Control is an illusion (0:00–0:12)

On screen while you're introduced. Let the dial drift.

0:00 Hi, I'm Kat. Every result in biology is a difference from a control, and nobody checks the control. So I'm building confidence without it. *Point at the pink dot, then the green ring. Click.*

02 Founder (0:12–0:40)

Self. Let the logos and the pink line under your name do the credential work.

0:12 Five labs, MGH, AbbVie, MIT, insitro, Axiom, and the same question at every bench: do my cells look normal? Nobody could answer it, and it made the science slow and wrong.

0:20 At Axiom we onboarded primary hepatocytes through a good CRO. Their checks passed; ours passed. We had noticed some uneven spreading in wells, and nothing could tell us whether it mattered. The first sign it did came from outside: the

Broad couldn't align our profiles to their reference. That's the gap. Not seeing something; knowing whether it matters.

0:35 This is the thing I kept needing and nobody sold. *Click.*

Founder detail, for the judge question (about 30 seconds). Three methods, three papers, one thread. At AbbVie I built a neuronal tau aggregation assay for drug discovery [3]: the readout was what the cells did, not a label on them. At Axiom I built Cell Painting in primary human hepatocytes for cytotoxicity and mode of action [1], ran assay development for two of the largest Cell Painting datasets in the world, and invented two bioactivity detection methods that see a small molecule's effect on cells at about ten times lower concentration than LDH, ATP or cell count. With the Carpenter lab we showed that cell count alone predicts small-molecule bioactivity benchmarks [2]: even the simplest phenotype carries the signal. The thread is that bioactivity is a measurement of behaviour, and behaviour is the half of QC nobody measures. Every lab already images its control wells; Datum Line turns that image into the certificate, against the field.

Precision to hold. The question predates plates: at MGH it was whether a *Xenopus* oocyte, a single frog egg cell, looked normal before a recording. Say only what your author contribution statements support: "built" on [1] and [3], "we showed" on [2]. The two bioactivity methods and the ten-fold sensitivity are from your resume (macrofluidic.me/resume); if asked, name the comparison assays and the concentration range. The 2019 GABA_A paper [5] is the MGH anesthesia lab, Harvard Medical School, 2017 to 2018. The industry claim is the deck's own: STR and vendor certificates cover identity, bioactivity covers behaviour, and behaviour is what a control well measures in any lab that images one.

Precision to hold. "Published with" lists coauthor institutions, never partners. The Axiom story has no villain: the CRO passed every check it had, Axiom passed every check it had, and the failure was that no check for this existed. Name no vendor and no person. The uneven distribution was noticed; its severity and its effect on data quality were unknown, so never present it as the cause, and never say "spun down

unevenly". If asked what it was: "We still don't know for certain. The uneven spreading is one hypothesis; whole-well images show it more easily than our fields, and a field reference is what would tell you whether it matters." The sequence is: local checks passed, uneven spreading noticed but unquantified, the Broad flagged that reference profiles would not align, the cause is still open. Say what happened, not how it was fixed. Confirm with Axiom that you can tell it; if not, say "at one lab".

03 Do my cells look normal? (0:40–1:05)

Us, then now. Point at the pink slice, then the two edges of card 3: run it, or hold.

0:40 This isn't my problem; it's all of ours. About twenty-eight billion dollars a year of US preclinical research can't be reproduced, and the biggest slice, about a third, is inputs.

0:52 Today you search, you ask around, then you guess. I have seen both edges. One lab ran it anyway: hundreds of thousands in R&D before the SOP was fixed. Another held: months developing cells that were already good enough. One too fast, one too slow. Cells are the input nobody certifies. Think growth chart: nobody plots their cells against the population. *Click.*

Precision to hold. \$28B is a 2015 US estimate; the 36% covers all reagents and reference materials, not only cells. The search card shows six real U2OS images from six papers, 2009 to 2023, via IDR and JUMP; the sources are in the Fig. 1 caption. Card 3 names neither lab and gives no exact figure; keep it that way on stage. The \$2.6M screen estimate stays on 06 and in the appendix. The why-now block on the slide cites JUMP [4]; the spoken why-now stays on 04.

04 One lab, or six? (1:05–1:45, three clicks)

Why now. State 1 is a game: let the room look for four seconds.

1:05 Six untreated wells. Same line, same protocol, no drug. One lab, or six? *Beat. Click.*

1:15 Six labs. Nobody here can tell; a classifier can, ninety-six percent of the time, against fourteen by chance. The pink one fools it. *Click.*

1:28 Why now: this data is public, eleven sites, free, and embeddings are a commodity. Same-size nuclei, different texture: that's what moves the point. The model isn't the hard part. The reference is. *Click.*

Precision to hold. Seven institutions, 12,046 wells, held out by plate; shuffled labels score 15.2%. Cell Painting fluorescence; brightfield is the bet.

05 It passed every check (1:45–2:20, three clicks)

How. Dark slide: slow down.

1:45 One plate. Everything its own lab can see says fine. *Click.*

1:52 It's labelled lab 2. Zero of its 384 untreated wells look like lab 2. *Click.*

2:02 And even one lab drifts: runs three months apart sit three times further apart than runs on the same day, then it plateaus. Your own past is not a reference. The field is. *Click.*

For ML people. The line is on the board in the second state of 04: batch correction takes lab identity from 96% to 12%, right for discovery, wrong for QC. Say it to the ML faces in the room.

06 Know on day one, not month four (2:20–2:45)

What. The sentence under the headline is the product; read it, then walk the four icons left to right. The certificate on the right is the real plate from 05, so point back at it.

2:20 The product: brightfield images of the control wells you already take. They stay yours. We place them against the field. One answer in a day: in or out, and which features moved.

2:35 That certificate is the plate you just saw: it passed every local check, and zero of its 384 wells look like its own lab. Four months becomes one day. *Click.*

07 Identity, yes. Behaviour, no. (2:45–2:55)

Point at the "run STR" cell. One sentence.

2:45 STR proves which line you have. Nothing proves it behaves like everyone else's, or says which way to move it back. STR is a collaborator, not a rival. *Click.*

08 Three buyers, one certificate (2:55–3:10)

Who buys. Walk the three columns; the pink row is what they pay.

2:55 Three buyers. Cores risk the paper; they contribute control images and use it free, and that builds the reference. Pharma and CROs risk the screen; they pay per cell line per quarter, about a hundred and seventy times less than one unchecked screen. Cell banks risk the lot; they pay per certified lot, and that builds the network. Proposed, not yet validated. *Click.*

Precision to hold. The prices are the README's proposal, not a validated model: free for contributors, ~\$5–15k per line per quarter for screening groups, a fee per certified lot. If you need ten seconds back, say only the first sentence of each column.

09 The ugliest dataset in biology (3:10–3:35)

The build. This is the blank-cheque answer, said before anyone asks. Point at the six real wells, then the four corrupted tiles.

3:10 If someone wrote the cheque, I'd build the ugliest dataset in biology, on purpose. These six wells are the same line, same protocol, no drug, and six labs already look like this: five microscope models, eight plate types. Then plant the failures: defocus, exposure, clipped wells, bad metadata.

3:23 Every method has to name which thing went wrong, not just that something did: a public benchmark, name the failure. A foundation other people can stand on, instead of Google Images. *Click.*

10 Every plate, placed in the field (3:35–4:00)

Vision, then the ask. Read the headline, then the right card.

3:35 That's the vision: every plate you image, placed in the field, without anyone lifting a finger. A growth chart for cell lines. Every certificate adds a plate to the reference, and the reference is the moat: nobody can build it alone. Next it tells you what to change, then flags a bad lot across every lab using it.

3:47 What I need: images, U2OS first; ML people better than me; bench time; and a pre-seed to run it. If you run a cell line somebody else also runs, you already have the data.

3:57 Have images of U2OS cells? Talk to me. *Stop.*

Precision to hold. Price and test-set budget are not set: say "per line, per quarter" and "the proof of concept sets the size". Don't name a partner lab until one has agreed.